

# Computer Vision 1: Exercise Sheet 2

Summary:

1. Read, display and write images using the packages `imageio` and `matplotlib`
2. Each image is represented by a multi-dimensional `numpy` array. Image manipulation is array manipulation.
3. Statistical features of images: mean, variance and histograms

## Read and display an image

The image in Figure 1 comes from the MSCOCO dataset<sup>1</sup>. Download it from Moodle with the name `sample.jpg`.



Figure 1: A sample image (id=345434) from the MSCOCO dataset

## Read an Image

The image can be read as follows:

```
1 import imageio
2 im = imageio.imread(uri='sample.jpg')
3 print(im.shape) # shape of the image
4 print(type(im)) # type of the variable that represents the image
5 print(im.dtype)
```

The output should be

```
1 (480, 640, 3)
2 <class 'imageio.core.util.Image'>
3 uint8
```

The shape of the array is (480,640,3) which corresponds to (Height, Width, Channels), with the three channels usually standing for red (R), green (G), and blue (B). The data type is **unsigned 8 bit integer** (what is the range of values for this data type?).

---

<sup>1</sup><http://cocodataset.org/>

The `type(im)` command does not give `<type 'numpy.ndarray'>` as we normally see for numpy arrays, but the class `imageio.core.util.Image` really is numpy multi-dimensional array in disguise<sup>2</sup>.

## Display an Image

One of the most popular Python packages for displaying images is `matplotlib`. An image can be visualized as follows:

```
1 import matplotlib.pyplot as plt
2 plt.imshow(im)
3 plt.show()
```

**Note:** Using `plt.show()` in a script will block the executing of the script until the figures have been closed.

A new window will pop out with our image displayed in it and some control buttons at the bottom-left corner, something like this:



Figure 2: Display a RGB image using `matplotlib`

Note that you can also see the numbering on the axes and verify the size of the image matches what you saw earlier.

To display a grayscale image, add `cmap='gray'` to the `imshow` function. Download the grayscale image `woman.png` from Moodle and display it:

```
1 im_grayscale = imageio.imread(uri='woman.png')
2 print(im.shape)
3 plt.imshow(im_grayscale, cmap='gray')
4 plt.show()
```

Expected output:

---

<sup>2</sup> If you have doubts about this, check out the source code at <https://github.com/imageio/imageio/blob/master/imageio/core/util.py>, line 130 and 191.



Figure 3: Display a grayscale image using `matplotlib`

`Matplotlib` is a very powerful visualization tool, but in this exercise we only use it to view images. After the exercise, you may want to explore what else `matplotlib` can do by following the online tutorial at <https://matplotlib.org/tutorials/index.html>.

- Repeat the steps above for another image in the MSCOCO dataset from <http://cocodataset.org/#explore>. Either load the image to disk before reading it, or directly read from the URL provided on the dataset webpage.

## Image Manipulation and writing images

Import `numpy` and set the seed of the random number generator (RNG) as follows:

```
1 import numpy as np
2 np.random.seed(0)
```

If you set a seed for the RNG, you will obtain the same sequence of pseudo-random numbers every time you run the next commands. This is useful for ensuring repeatability of experiments.

### Adding noise

Use `np.random.normal` to create normally distributed noise. The documentation can be found at <https://docs.scipy.org/doc/numpy-1.15.0/reference/generated/numpy.random.normal.html>. Add the created noise array to the image array, and display the image. Note the following:

- The shape of the noise array must match the shape of the image array.
- The created noise has a data type of `float64`. When you add it to an image with type `uint8`, the result will have the type `float64`. Before displaying the image, you must convert the result back to `uint8` by using Numpy's `astype` method. To avoid overflow



and underflow in the conversion, use the `clip` method to ensure the values are in the appropriate range for `uint8` (what are the minimum and maximum value representable as an unsigned 8-bit integer?).

- Try different inputs for the mean and standard deviation of the Gaussian, these are represented by the input parameters `loc` and `scale` of `np.random.normal`. How do they visually affect the image?

An example output with `loc=0.0` and `scale=20.0` is shown in Figure . Can you see how it differs from Figure 1?



Figure 4: Original image.



Figure 5: Noisy image.

## Flipping and rotating an image

- Use `np.fliplr` and `np.flipud` to flip the image, and display the result.
- Use `np.rot90` to rotate the image and display the result.

Other less basic operations such as resizing, cropping and rotating with arbitrary angles are not very well supported by `numpy`. You can use other modules and libraries such as `OpenCV` and `Skimage` to do this.

## Write an image to file

Write any modified image from above to a file using `imageio`. You should use the function `imwrite`, the documentation can be found at <https://imageio.readthedocs.io/en/stable/userapi.html#imageio.imwrite> – you should only need the `uri` and `im` input arguments. Check that you can find the image on disk, and that it looks the same as your visualization with `matplotlib`.

## Statistical features: Mean, variance and histogram

`Numpy` provides built-in functions for computing statistical image features such as the mean, variance, and histogram.

- Find the documentation of `np.mean`, `np.var`, and `np.histogram`.

- Calculate the mean and variance of the R,G,B channels of an image using just one call to `np.mean` and `np.var`. The expected outputs should have a shape `(3,)`. Hint: use the `axis` argument to calculate the mean and variance along the width and height dimensions only.
- Compute and display the histograms of each of the three channels in the image: red, green, and blue. Follow the examples on the documentation page of `np.histogram`. How does the result change if you use a different number of bins? What if you manually specify which bins to use?

Hints:

- `np.histogram` flattens the input array, so only give the pixel values in one channel (R, G, or B) at a time. Do you remember from the previous exercise how to index Numpy arrays?
- Take two output arguments from `np.histogram`:

```
1 hist, bin_edges = np.histogram(...)
```

Then calculate the display width and create a Numpy array for the histogram bin centers as follows:

```
1 width = 0.8 * (bin_edges[1] - bin_edges[0])
2 bin_centers = (bin_edges[:-1] + bin_edges[1:]) / 2
```

Then you can draw the histogram using `plt.bar`:

```
1 plt.bar(bin_centers, hist, width=width)
```

If you want, you can add a title to a plot, e.g., by `plt.title('Red')`. Use `plt.xlabel()` and `plt.ylabel()` for axis labels.

## Histogram in a bounding box

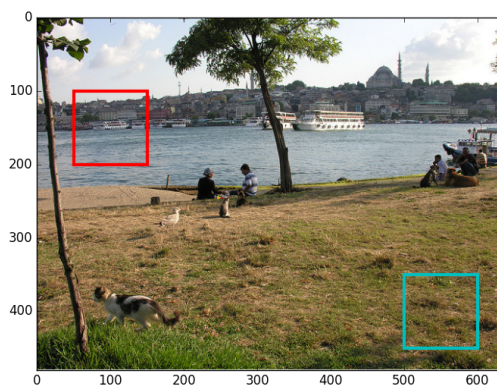


Figure 6: Figure with bounding boxes.

Figure 6 shows the image from the exercise with two bounding boxes, one red and one cyan. The upper left corner of the red bounding box is at `(50, 100)`, and the upper left corner of the cyan bounding box is at `(500, 350)`. Both bounding boxes have a size of 100-by-100 pixels. Load the sample image as `uint8`.

- For the pixels inside the red bounding box, calculate and draw the R, G, and B channel histograms for the using 20 equally spaced bins in the range between 0 and 255.
  - Use `numpy.linspace` to create an equally spaced vector to use as bins.
- Repeat the same for the pixels inside the cyan bounding box.
- Calculate the average RGB value inside each of the two bounding boxes.